


Matrix Multiplication Analysis




Nodebechukwu Okoye | Jakob Lopez

Introduction

There are a lot of ways to perform matrix multiplication

Examples are Iterative algorithm, Divide and conquer algorithm, Non-square matrices, Sub-cubic algorithms, Parallel and distributed algorithms.



Naive Algorithm

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \times \begin{bmatrix} e & f \\ g & h \end{bmatrix} = \begin{bmatrix} ae + bg & af + bh \\ ce + dg & cf + dh \end{bmatrix}$$

A B C

A, B and C are square matrices of size $N \times N$

a, b, c and d are submatrices of A, of size $N/2 \times N/2$

e, f, g and h are submatrices of B, of size $N/2 \times N/2$

Strassen's Algorithm

$$\begin{aligned}p_1 &= a(f - h) \\p_3 &= (c + d)e \\p_5 &= (a + d)(e + h) \\p_7 &= (a - c)(e + f)\end{aligned}$$

$$\begin{aligned}p_2 &= (a + b)h \\p_4 &= d(g - e) \\p_6 &= (b - d)(g + h)\end{aligned}$$

The A x B can be calculated using above seven multiplications.
Following are values of four sub-matrices of result C

$$\begin{array}{c} \left[\begin{array}{c|c} a & b \\ \hline c & d \end{array} \right] \times \left[\begin{array}{c|c} e & f \\ \hline g & h \end{array} \right] = \left[\begin{array}{c|c} p_5 + p_4 - p_2 + p_6 & p_1 + p_2 \\ \hline p_3 + p_4 & p_1 + p_5 - p_3 - p_7 \end{array} \right] \\ A \qquad \qquad B \qquad \qquad \qquad C \end{array}$$

A, B and C are square matrices of size $N \times N$

a, b, c and d are submatrices of A, of size $N/2 \times N/2$

e, f, g and h are submatrices of B, of size $N/2 \times N/2$

p1, p2, p3, p4, p5, p6 and p7 are submatrices of size $N/2 \times N/2$

Design Considerations

- Matrix sizes
 - $2048 \times 2048 (2^{11} \times 2^{11})$
 - $8192 \times 8192 (2^{13} \times 2^{13})$
 - $16384 \times 16384 (2^{14} \times 2^{14})$
- Processes (Strassen)
 - 1 for Serial
 - 7 for Parallel
- Threads (Naive)
 - 1 for Serial
 - 4 for Parallel (Control)
 - 8 for Parallel
 - 16 for Parallel

Stampede2 Hardware Specs

- Processor: Intel Xeon Phi 7250 ("Knights Landing")
- Cores per node: 68 cores on a single socket
- Threads per core: 4
- Threads per node: 272
- Clock rate: 1.4 GHz
- RAM: 96GB DDR4 plus 16GB high-speed MCDRAM
- Cache: 32KB L1 data cache per core; 1MB L2 per two-core tile
- 18 petaflop

Software and Tools

Python

Mpi4py : for MPI

Numpy : for all matrix calculation

Multiprocessing : for threads

Time: for timing the code

Stampede 2 : For the cluster

Collaboration

Git and Github

Compilation and Execution Commands

- Submission
 - Sbatch <script name>

```
#
# -- If you're running out of memory, try running
#    fewer tasks per node to give each task more memory.
#
#-----

#SBATCH -J OkoyeLopezProjectNaiveParallel      # Job name
#SBATCH -o 16_16384.o%j                        # Name of stdout output file
#SBATCH -e OkoyeLopezProjectNaiveParallel.e%j  # Name of stderr error file
#SBATCH -p normal                             # Queue (partition) name
#SBATCH -N 1                                  # Total # of nodes
#SBATCH -n 1                                  # Total # of mpi tasks
#SBATCH -t 45:30:00                           # Run time (hh:mm:ss)
#SBATCH --mail-user=
#SBATCH --mail-type=all                       # Send email at begin and end of job
#SBATCH -A CMPS-5433-MWSU                     # class project/account code

# Other commands must follow all #SBATCH directives...

module list
pwd
date

# Launch MPI code...
module spider python3/3.7.0
python OkoyeLopezProjectNaiveParallel.py
# -----
```


Execution Time(seconds): Naive

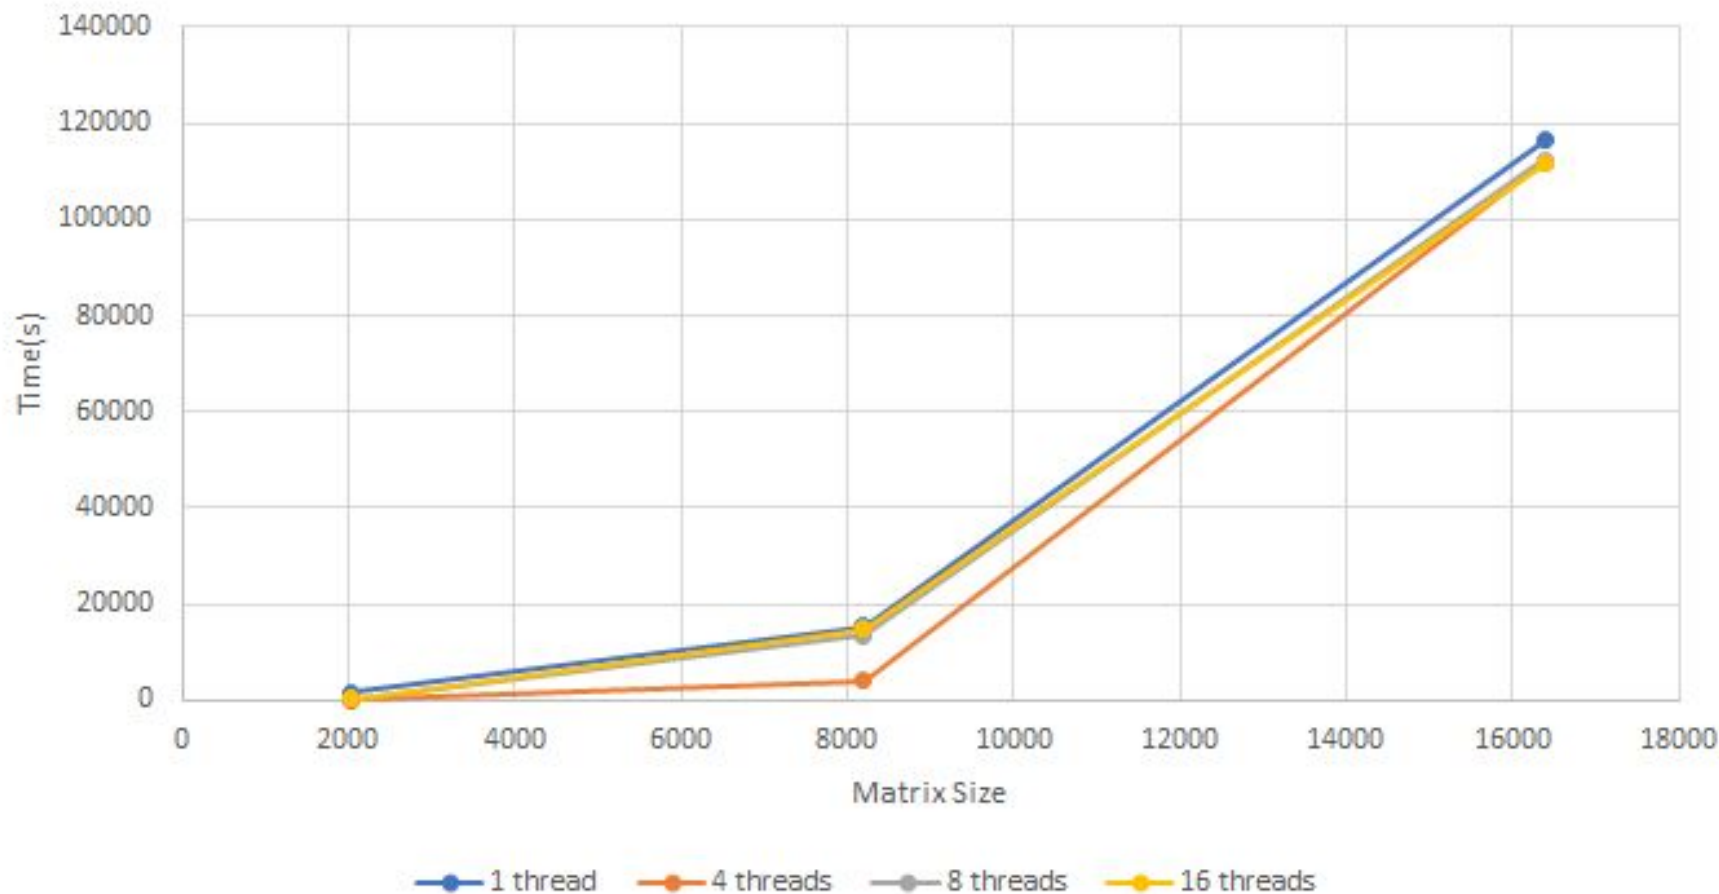
Matrix Size

Threads

	2048	8192	16,384
1	1722	15,299	116,464
4	68	4047	111,989
8	225	13,415	112,378
16	258	14,712	111,580

* 1 day = 86,400 seconds

Naive Algorithm Execution Time



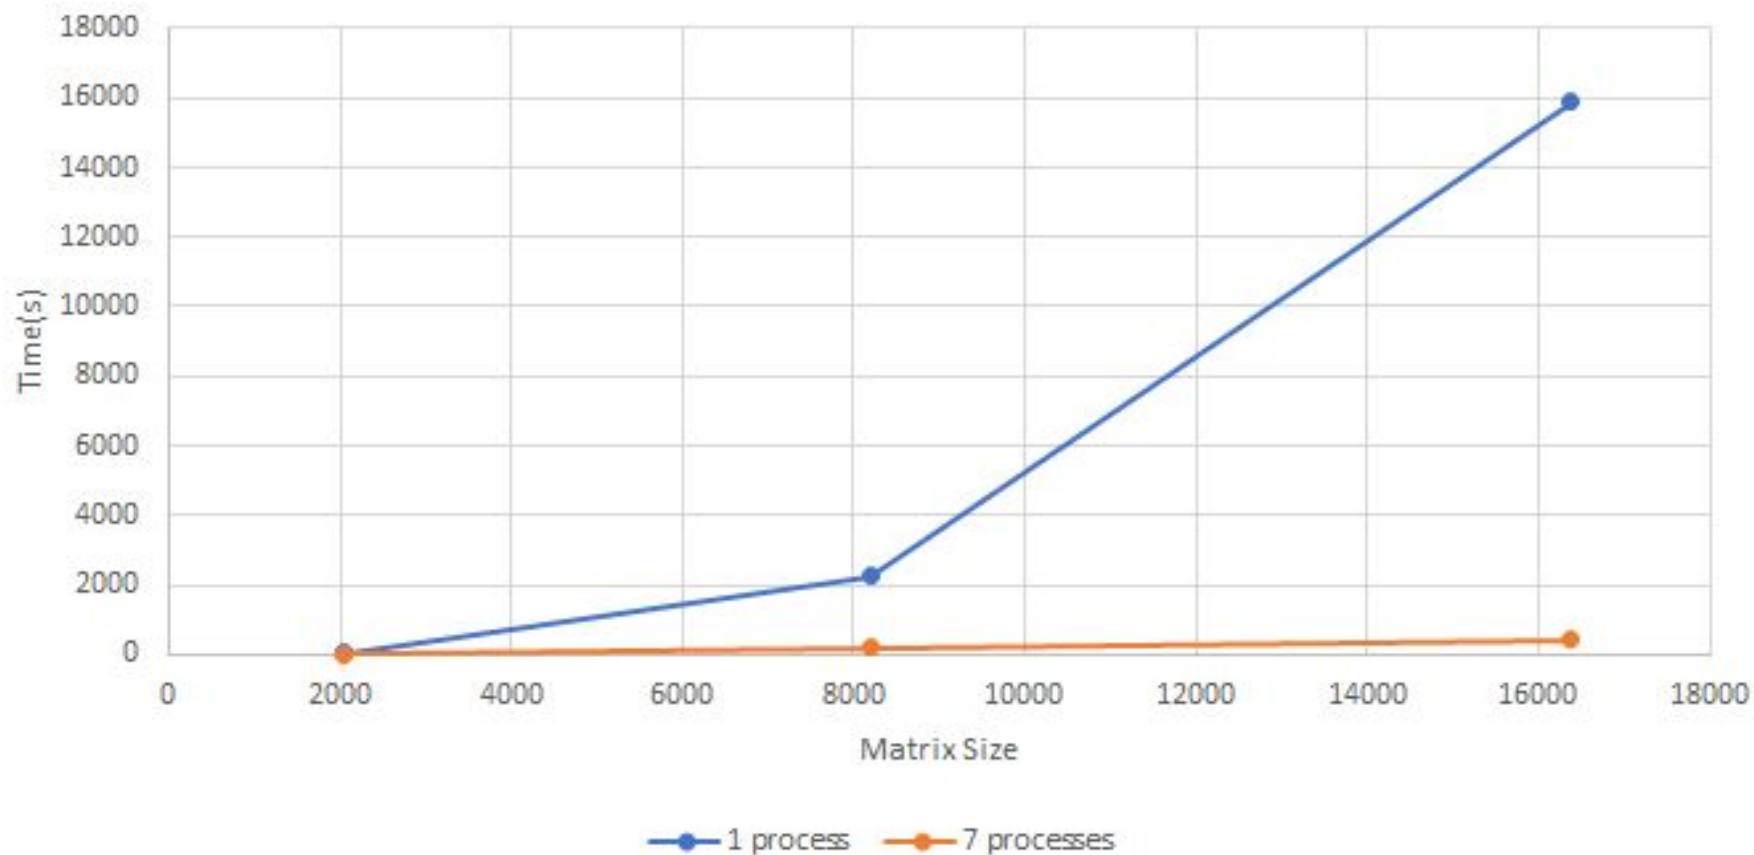
Execution Time(seconds): Strassen

Processes

Matrix Size

	2048	8192	16,384
1	45	2266	15,872
7	4	202	1407

Strassen's Algorithm Execution time



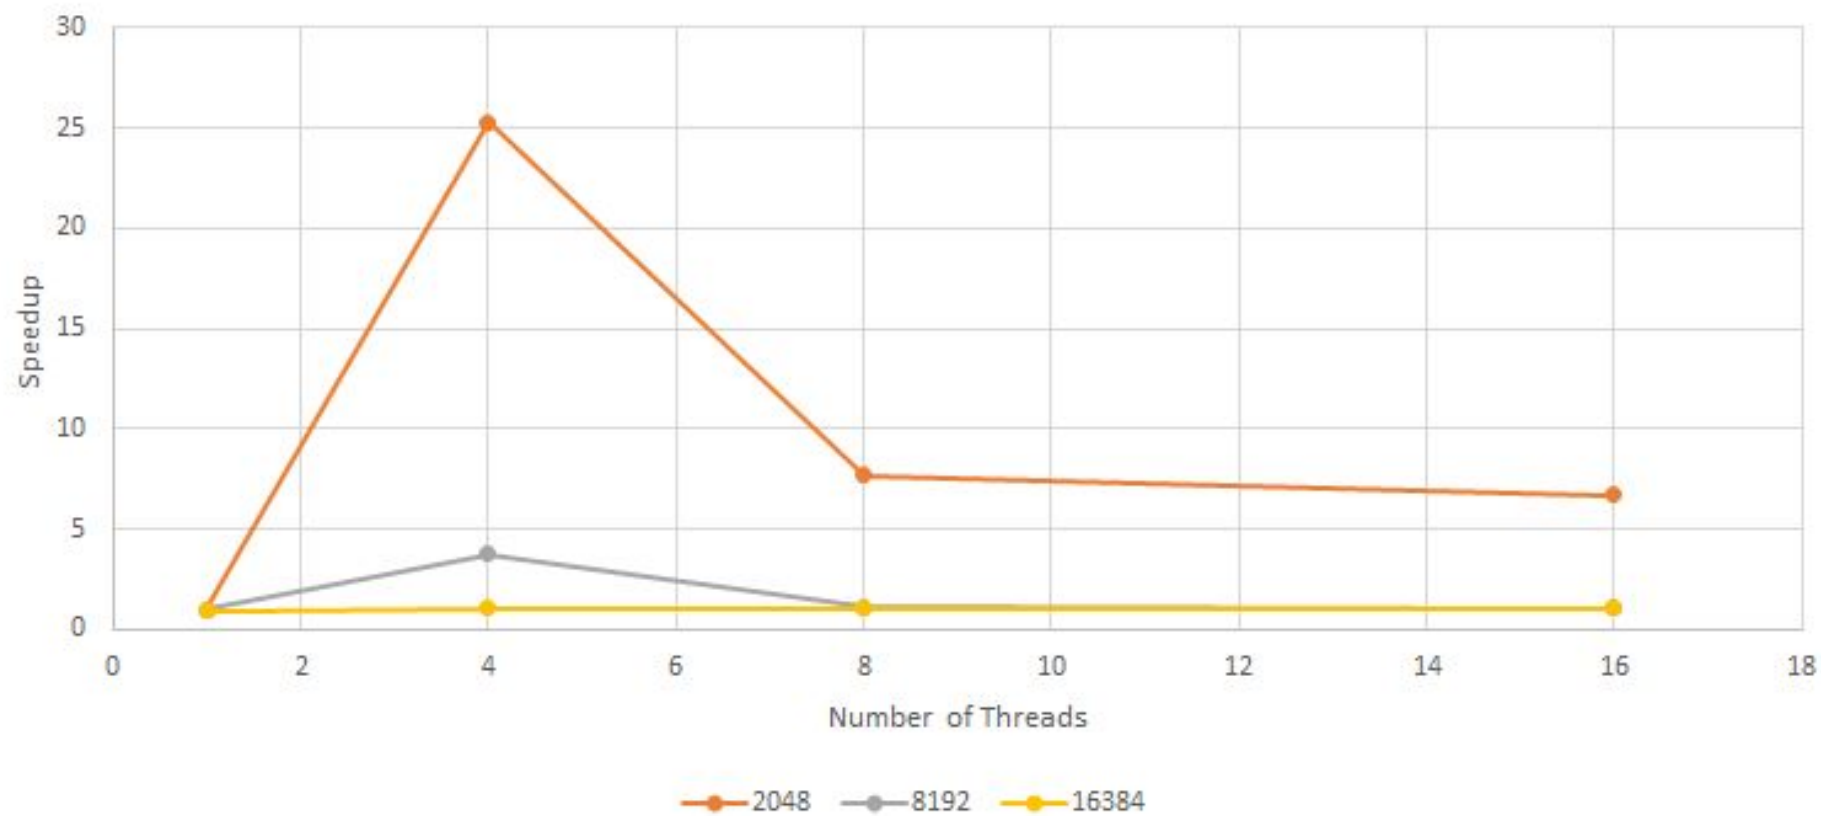
Speedup: Naive

Matrix Size

Threads

	2048	8192	16,384
1	1	1	1
4	25.32353	3.780331	1.039959
8	7.653333	1.14044	1.036359
6	6.674419	1.039899	1.043771

Naive Speedup



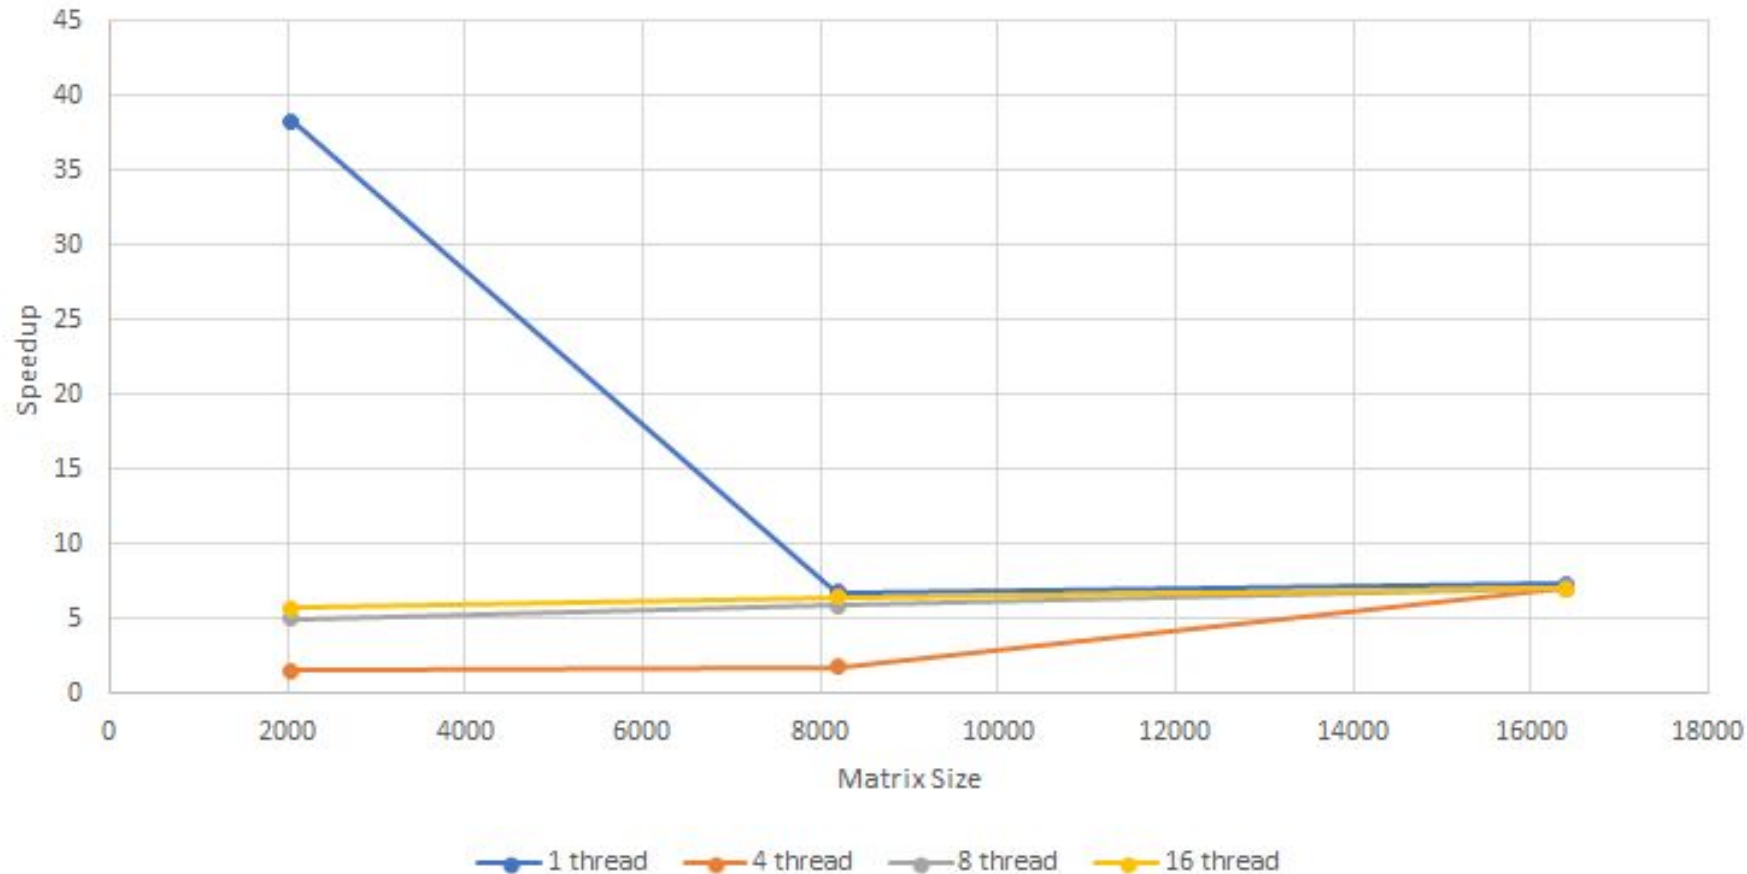
Speedup: Serial Strassen

Matrix Size

Naive Threads

	2048	8192	16,384
1	38.26667	6.751544572	7.337702
4	1.511111	1.785966461	7.055759
8	5	5.920123566	7.080267
16	5.733333	6.492497793	7.02999

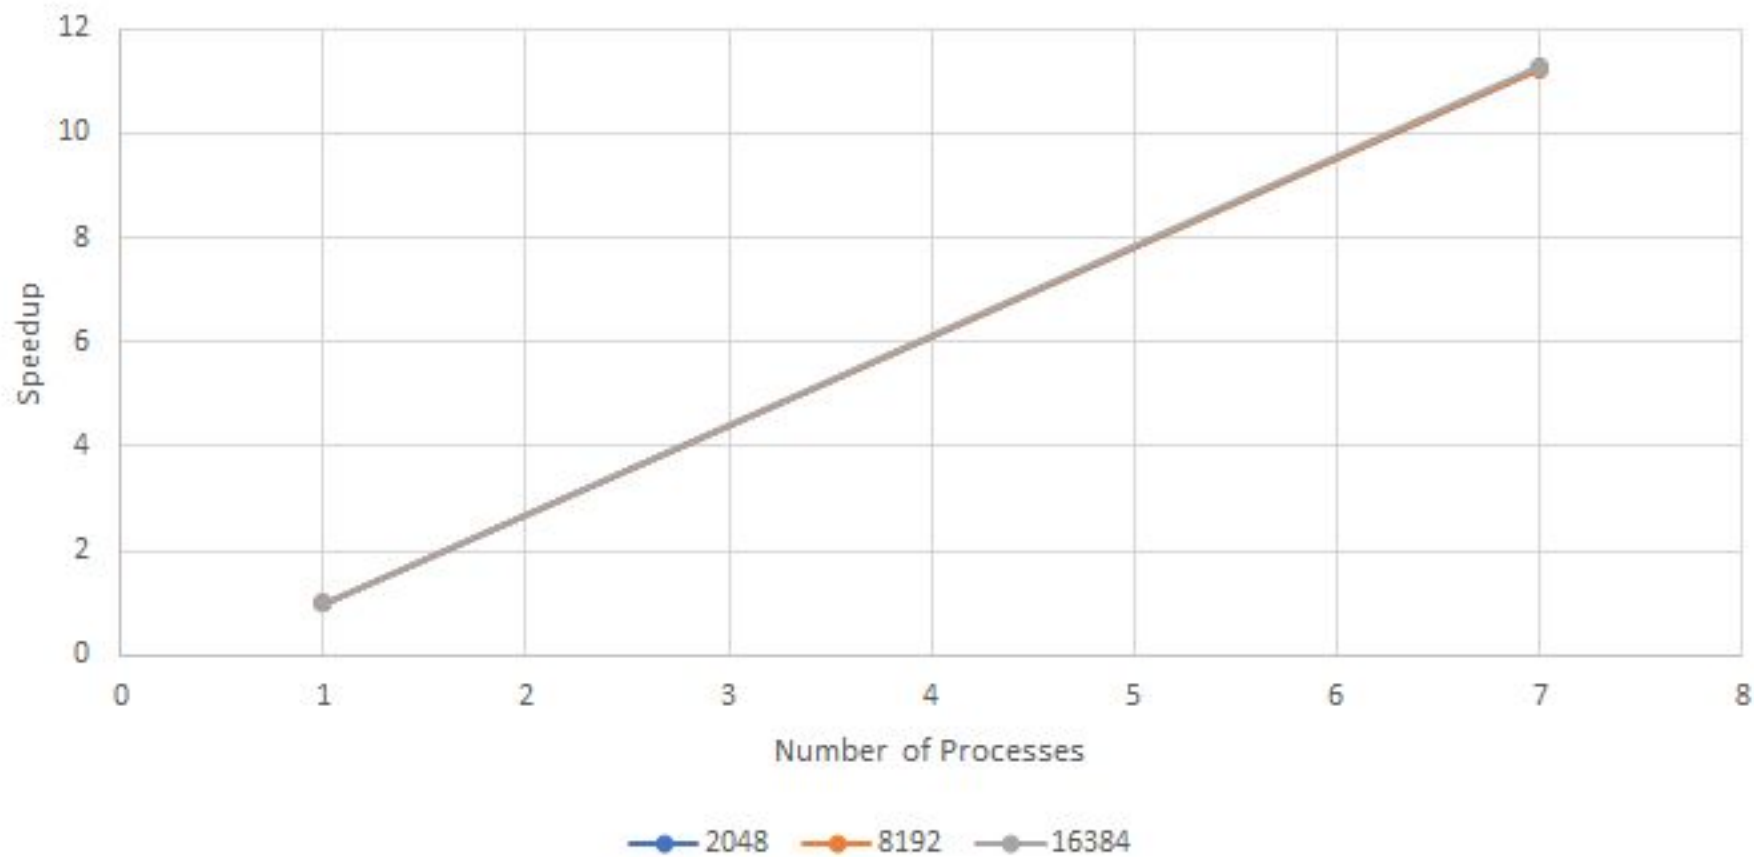
Algorithm Speedup: Serial Strassen vs Naive



Speedup: Strassen

Processes	Matrix Size			
		2048	8192	16,384
	1	1	1	1
	7	11.25	11.21782	11.28074

Strassen Speedup



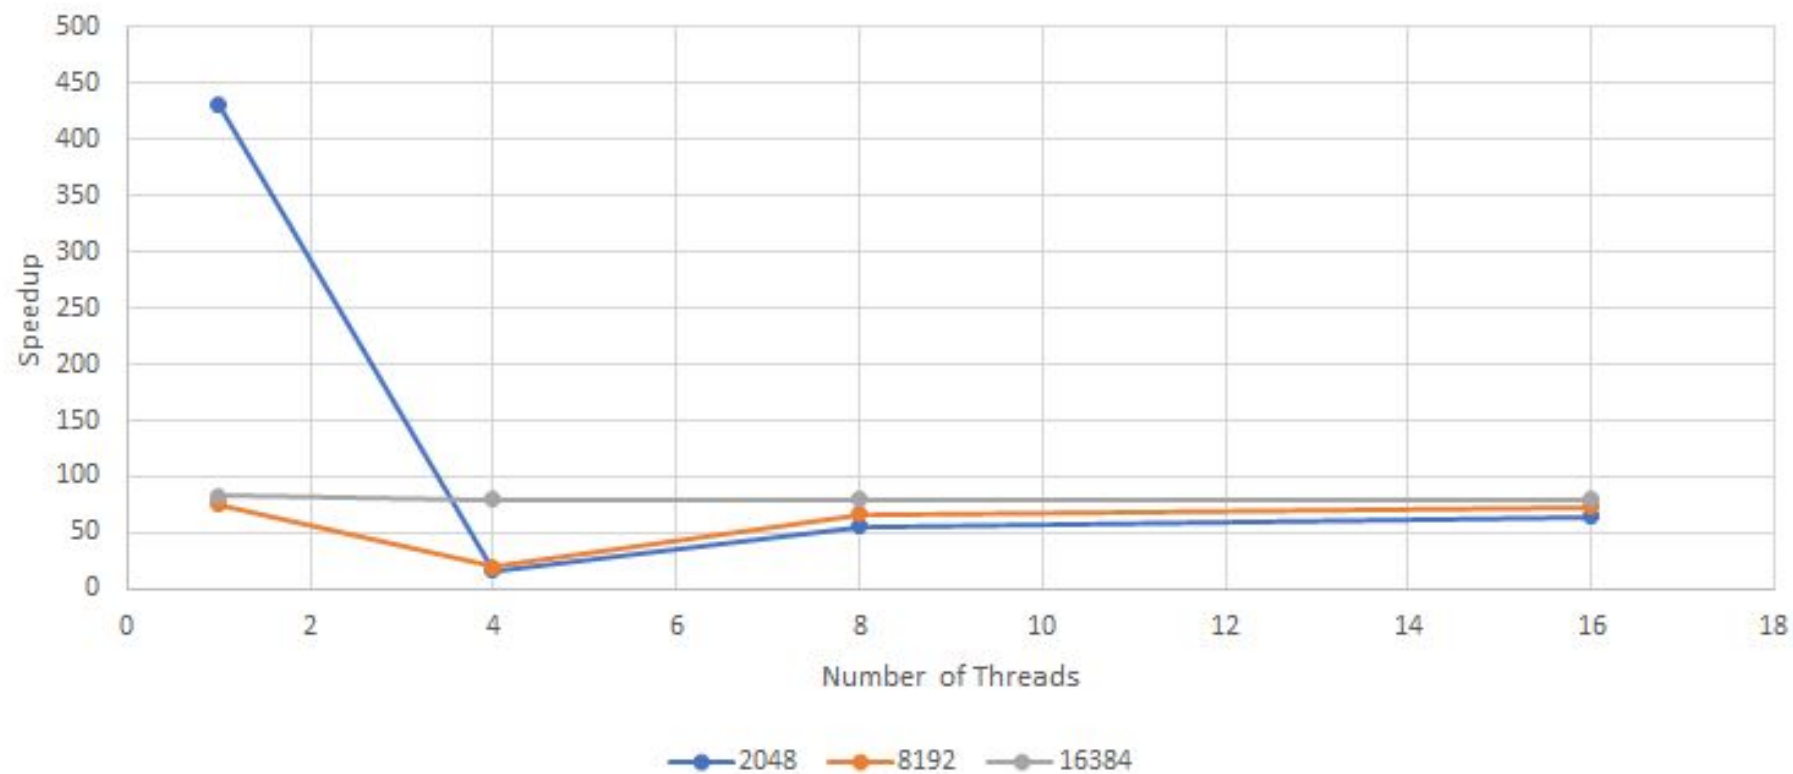
Speedup: Parallel Strassen

Matrix Size

Naive Threads

	2048	8192	16,384
1	430.5	75.73762	82.7747
4	17	20.03465	79.59417
8	56.25	66.41089	79.87065
16	64.5	72.83168	79.30348

Algorithm Speedup: Parallel Strassen vs Naive

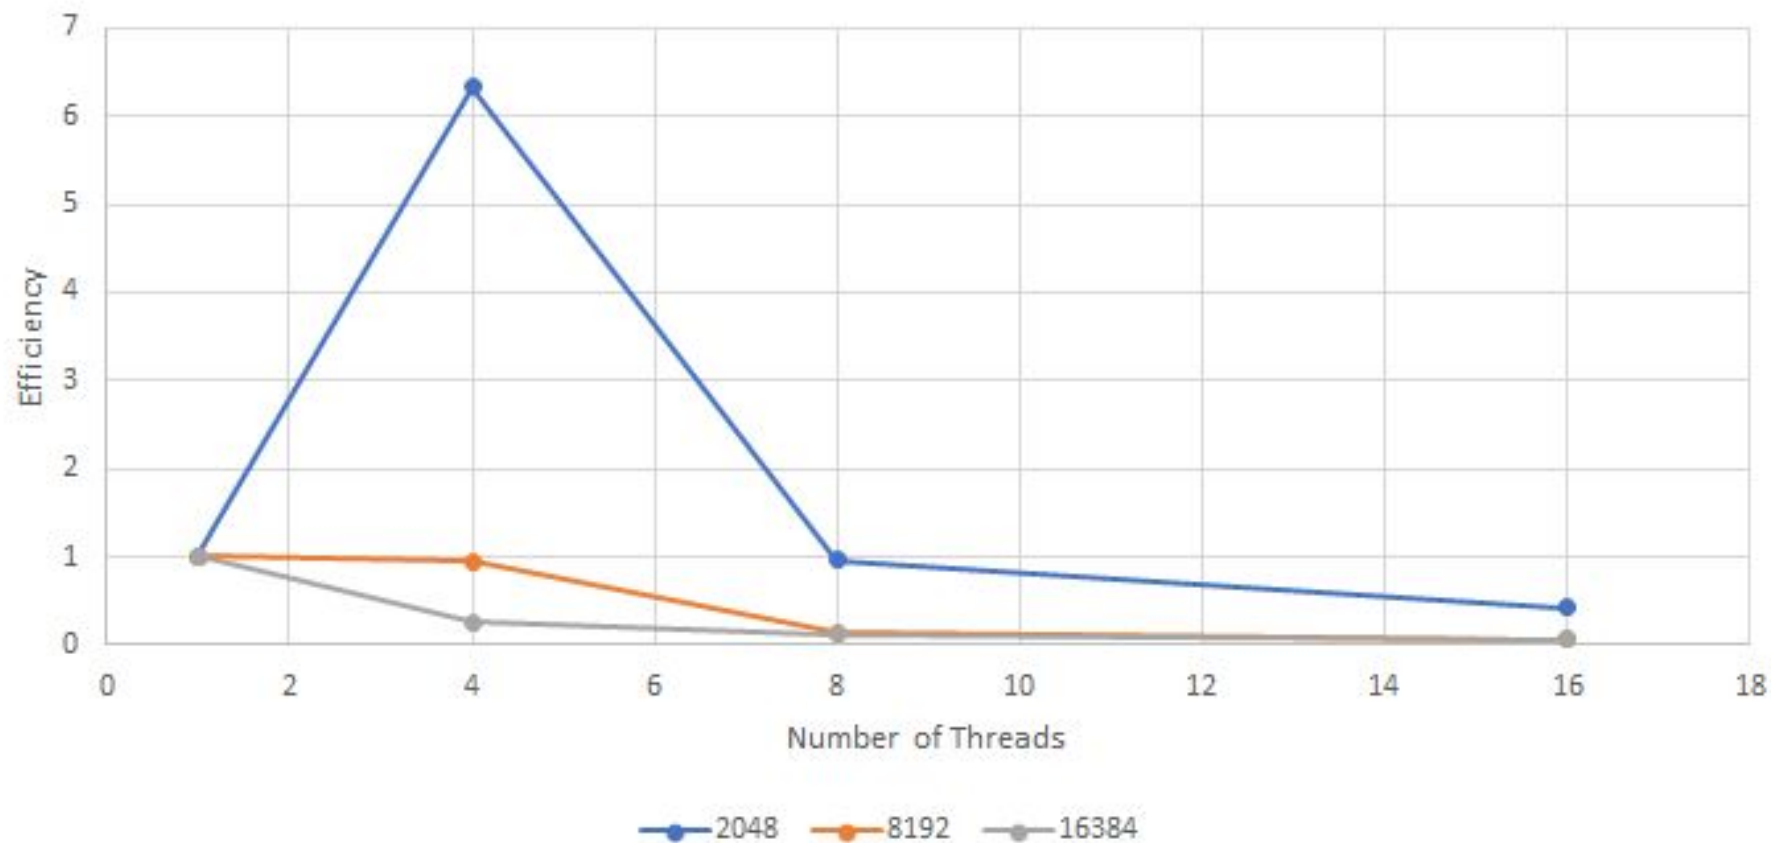


Efficiency: Naive

Matrix Size

Threads	Matrix Size			
		2048	8192	16,384
	1	1	1	1
	4	6.330882	0.945083	0.25999
	8	0.956667	0.142555	0.129545
	16	0.417151	0.064994	0.065236

Naive Efficiency




Efficiency: Strassen

Processes

Matrix Size

	2048	8192	16,384
1	1	1	1
7	1.607143	1.602546	1.611534

Strassen Efficiency




Efficiency: Parallel Strassen

Matrix Size

Naive Threads

	2048	8192	16,384
1	430.5	75.74	82.7747
4	4.25	5.009	19.89854
8	7.03125	8.301	9.983831
16	4.03125	4.552	4.956468

Algorithm Efficiency: Parallel Strassen vs Naive



Important Variables


- Size of matrix
 - `n`
- Number of threads
 - `num_workers`
- Number of processes
 - Change in Parallel Strassen Script, however not advised

Obstacles

- New language for us to use in Stampede2
- Understanding Strassen's Algorithm
- Massive waiting time
- Inefficiency with Microsoft Excel

Conclude

- Naive algorithm is not the only matrix multiplication algorithm
- Strassen's algorithm is the most practical big matrix multiplication technique
- Parallel processing is made easy by Python



<https://github.com/JakobLopez/ParallelMatrixMultiplication>

References


https://www.cac.cornell.edu/education/training/StampedeJan2017/Python_R_HPC_Workshop.pdf

<https://mpi4py.readthedocs.io/en/stable/>

<https://docs.python.org/2/library/multiprocessing.html>

https://en.wikipedia.org/wiki/Matrix_multiplication_algorithm

<https://www.geeksforgeeks.org/strassens-matrix-multiplication/>

 naive_parallel.py

```
1  if questions:
2      ask()
3  else:
4      clap()
```